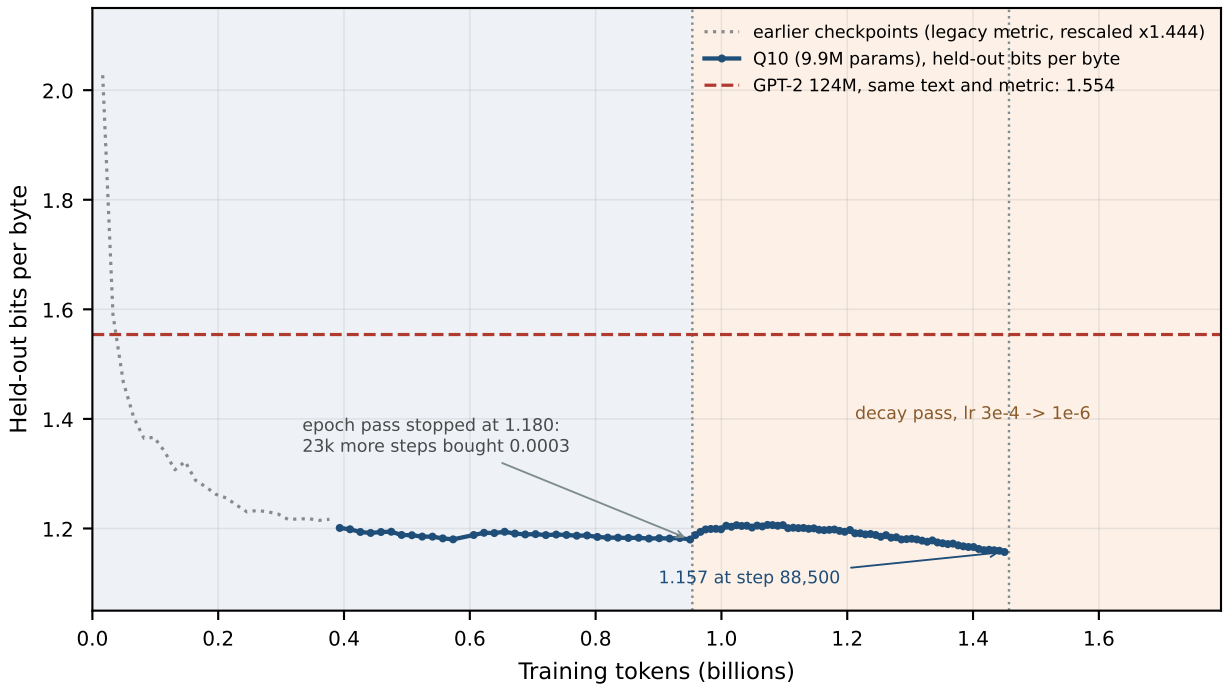
earlier checkpoints (legacy metric, rescaled x1.444)
Q10 (9.9M params), held-out bits per byte
GPT-2 124M, same text and metric: 1.554
epoch pass stopped at 1.180:
23k more steps bought 0.0003
decay pass, lr 3e-4 -> 1e-6
1.157 at step 88,500
Held-out bits per byte
Training tokens (billions)
2.0
1.8
1.6
1.4
1.2
0.0
0.2
0.4
0.6
0.8
1.0
1.2
1.4
1.6